

Implementing a Neural Network in Python

WildML

March 28, 2016

Outline

- 1 How to install Python
 - Python and IDEs
 - Libraries
 - Anaconda Platform
- 2 Neural Network
 - Recap: Neural Networks
 - How our network makes predictions
 - Learning the Parameters
- 3 Neural Network in Python
 - Generating a dataset
 - Defining variables
 - Loss function
 - Predict function
 - Batch gradient descent

Outline

1 How to install Python

- Python and IDEs

- Libraries

- Anaconda Platform

2 Neural Network

- Recap: Neural Networks

- How our network makes predictions

- Learning the Parameters

3 Neural Network in Python

- Generating a dataset

- Defining variables

- Loss function

- Predict function

- Batch gradient descent

Python and IDEs

Install Python and pip

■ Linux (Ubuntu):

```
1 $ sudo apt-get install python
2 $ sudo apt-get install python-pip
3 $ sudo pip install --upgrade pip
```

■ Windows: Read [this article](#).

■ MacOS: Read [this article](#).

Python and IDEs

Python IDEs

- PyCharm
- WingIDE
- PyDev
- Vim and [spf-13](#)

Outline

1 How to install Python

- Python and IDEs
- Libraries
- Anaconda Platform

2 Neural Network

- Recap: Neural Networks
- How our network makes predictions
- Learning the Parameters

3 Neural Network in Python

- Generating a dataset
- Defining variables
- Loss function
- Predict function
- Batch gradient descent

- Download and install: Read [this article](#).

- Download and install: Read [this article](#).

Outline

1 How to install Python

- Python and IDEs
- Libraries
- **Anaconda Platform**

2 Neural Network

- Recap: Neural Networks
- How our network makes predictions
- Learning the Parameters

3 Neural Network in Python

- Generating a dataset
- Defining variables
- Loss function
- Predict function
- Batch gradient descent

Anaconda Platform

<https://www.continuum.io/>

- Download and install: Read [this article](#).

Outline

- 1 How to install Python
 - Python and IDEs
 - Libraries
 - Anaconda Platform
- 2 Neural Network
 - **Recap: Neural Networks**
 - How our network makes predictions
 - Learning the Parameters
- 3 Neural Network in Python
 - Generating a dataset
 - Defining variables
 - Loss function
 - Predict function
 - Batch gradient descent

Recap: Neural Networks




Figure : A 3-layer neural network with one input layer, one hidden layer, and one output layer.

Recap: Neural Networks

Cont.

- The number of nodes in the input layer is determined by the dimensionality of input data.
- The number of nodes in the output layer is determined by the number of classes.
- The number of nodes in the hidden layer is **hyper-parameter**, we can choose in the experiments.

Outline

- 1 How to install Python
 - Python and IDEs
 - Libraries
 - Anaconda Platform
- 2 Neural Network
 - Recap: Neural Networks
 - **How our network makes predictions**
 - Learning the Parameters
- 3 Neural Network in Python
 - Generating a dataset
 - Defining variables
 - Loss function
 - Predict function
 - Batch gradient descent

How our network makes predictions

Our network makes predictions using forward propagation:

$$z_1 = xW_1 + b_1$$

$$a_1 = \tanh(z_1)$$

$$z_2 = a_1W_2 + b_2$$

$$a_2 = \hat{y} = \text{softmax}(z_2)$$

where

- z_i is the input of layer i
- a_i is the output of layer i after applying the activation function
- W_1, b_1, W_2, b_2 are parameters of our network, which we need to learn from our training data

How our network makes predictions

Cont.

- Some activation functions often use:

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

$$\text{sigmoid}(x) = \frac{1}{1 + e^{-x}}$$

$$\text{relu}(x) = \max(0, x)$$

- *softmax* function is simply a way to convert raw scores to probabilities.

$$\text{softmax}(\mathbf{x})_j = \frac{e^{x_j}}{\sum_{k=1}^K e^{x_k}}$$

Outline

- 1 How to install Python
 - Python and IDEs
 - Libraries
 - Anaconda Platform
- 2 Neural Network
 - Recap: Neural Networks
 - How our network makes predictions
 - Learning the Parameters
- 3 Neural Network in Python
 - Generating a dataset
 - Defining variables
 - Loss function
 - Predict function
 - Batch gradient descent

Learning the Parameters

- Learning the parameters for our network means finding parameters (W_1, b_1, W_2, b_2) that minimize the error on our training data.
- A common choice with the softmax output is the **categorical cross-entropy loss**

$$L(y, \hat{y}) = -\frac{1}{N} \sum_{n \in N} \sum_{i \in C} y_{n,i} \log \hat{y}_{n,i}$$

Learning the Parameters

Cont.

- We can use **gradient descent** to find the minimum.
- Gradient descent needs the gradients (vector of derivatives) of the loss function with respect to our parameters: $\frac{\partial L}{\partial W_1}$, $\frac{\partial L}{\partial b_1}$, $\frac{\partial L}{\partial W_2}$, $\frac{\partial L}{\partial b_2}$
- To calculate these gradients we use the famous **back-propagation algorithm**.

Learning the Parameters

Cont.

Applying the back-propagation formula we find the following:

$$\delta_3 = \hat{y} - y$$

$$\delta_2 = (1 - \tanh^2 z_1) \circ \delta_3 W_2^T$$

$$\frac{\partial L}{\partial W_2} = a_1^T \delta_3$$

$$\frac{\partial L}{\partial b_2} = \delta_3$$

$$\frac{\partial L}{\partial W_1} = x^T \delta_2$$

$$\frac{\partial L}{\partial b_1} = \delta_2$$

Outline

- 1 How to install Python
 - Python and IDEs
 - Libraries
 - Anaconda Platform
- 2 Neural Network
 - Recap: Neural Networks
 - How our network makes predictions
 - Learning the Parameters
- 3 Neural Network in Python
 - Generating a dataset
 - Defining variables
 - Loss function
 - Predict function
 - Batch gradient descent

Implement

Generating a dataset

Example

scikit-learn has some useful dataset generators, so we don't need to write the code ourselves:

```
1 import matplotlib.pyplot as plt
2 import numpy as np
3 import sklearn
4 import sklearn.datasets
5 import sklearn.linear_model
6 import matplotlib
7
8 # Generate a dataset and plot it
9 np.random.seed(0)
10 X, y = sklearn.datasets.make_moons(200, noise=0.20)
11 plt.scatter(X[:,0], X[:,1], s=40, c=y, cmap=plt.cm.
    Spectral)
```

Outline

- 1 How to install Python
 - Python and IDEs
 - Libraries
 - Anaconda Platform
- 2 Neural Network
 - Recap: Neural Networks
 - How our network makes predictions
 - Learning the Parameters
- 3 Neural Network in Python
 - Generating a dataset
 - Defining variables
 - Loss function
 - Predict function
 - Batch gradient descent

Implement

Defining variables

We start by defining some useful variables and parameters for gradient descent:

```
1 num_examples = len(X) # training set size
2 nn_input_dim = 2 # input layer dimensionality
3 nn_output_dim = 2 # output layer dimensionality
4
5 # Gradient descent parameters (I picked these by hand)
6 epsilon = 0.01 # learning rate for gradient descent
7 reg_lambda = 0.01 # regularization strength
```


Outline

- 1 How to install Python
 - Python and IDEs
 - Libraries
 - Anaconda Platform
- 2 Neural Network
 - Recap: Neural Networks
 - How our network makes predictions
 - Learning the Parameters
- 3 Neural Network in Python
 - Generating a dataset
 - Defining variables
 - **Loss function**
 - Predict function
 - Batch gradient descent

Implement

Loss function

First let's implement the loss function we defined above:

```
1 def calculate_loss(model):
2     W1, b1, W2, b2 = model['W1'], model['b1'], model['
3     W2'], model['b2']
4     # Forward propagation to calculate our predictions
5     z1 = X.dot(W1) + b1
6     a1 = np.tanh(z1)
7     z2 = a1.dot(W2) + b2
8     exp_scores = np.exp(z2)
9     probs = exp_scores / np.sum(exp_scores, axis=1,
10    keepdims=True)
11    # Calculating the loss
12    correct_logprobs = -np.log(probs[range(num_examples)
13    , y])
14    data_loss = np.sum(correct_logprobs)
15    # Add regularization term to loss (optional)
16    data_loss += reg_lambda/2 * (np.sum(np.square(W1))
17    + np.sum(np.square(W2)))
18    return 1./num_examples * data_loss
```

Outline

- 1 How to install Python
 - Python and IDEs
 - Libraries
 - Anaconda Platform
- 2 Neural Network
 - Recap: Neural Networks
 - How our network makes predictions
 - Learning the Parameters
- 3 Neural Network in Python
 - Generating a dataset
 - Defining variables
 - Loss function
 - Predict function
 - Batch gradient descent

Implement

Predict function

We also implement a helper function to calculate the output of the network

```
1 # Helper function to predict an output (0 or 1)
2 def predict(model, x):
3     W1, b1, W2, b2 = model['W1'], model['b1'], model['
4     W2'], model['b2']
5     # Forward propagation
6     z1 = x.dot(W1) + b1
7     a1 = np.tanh(z1)
8     z2 = a1.dot(W2) + b2
9     exp_scores = np.exp(z2)
10    probs = exp_scores / np.sum(exp_scores, axis=1,
        keepdims=True)
11    return np.argmax(probs, axis=1)
```

Outline

- 1 How to install Python
 - Python and IDEs
 - Libraries
 - Anaconda Platform
- 2 Neural Network
 - Recap: Neural Networks
 - How our network makes predictions
 - Learning the Parameters
- 3 Neural Network in Python
 - Generating a dataset
 - Defining variables
 - Loss function
 - Predict function
 - Batch gradient descent

Implement

Batch gradient descent

Finally, here comes the function to train our Neural Network.

```
1 def build_model(nn_hdim, num_passes=20000, print_loss=False):  
2     # Initialize the parameters to random values. We  
3     # need to learn these.  
4     np.random.seed(0)  
5     W1 = np.random.randn(nn_input_dim, nn_hdim) / np.  
6     sqrt(nn_input_dim)  
7     b1 = np.zeros((1, nn_hdim))  
8     W2 = np.random.randn(nn_hdim, nn_output_dim) / np.  
9     sqrt(nn_hdim)  
10    b2 = np.zeros((1, nn_output_dim))  
11    # This is what we return at the end  
12    model = {}
```

Implement

Batch gradient descent (Cont.)

Finally, here comes the function to train our Neural Network.

```
1 # Gradient descent. For each batch ...
2 for i in xrange(0, num_passes):
3     # Forward propagation
4     z1 = X.dot(W1) + b1
5     a1 = np.tanh(z1)
6     z2 = a1.dot(W2) + b2
7     exp_scores = np.exp(z2)
8     probs = exp_scores / np.sum(exp_scores, axis=1,
9                                   keepdims=True)
10    # Backpropagation
11    delta3 = probs
12    delta3[range(num_examples), y] -= 1
13    dW2 = (a1.T).dot(delta3)
14    db2 = np.sum(delta3, axis=0, keepdims=True)
15    delta2 = delta3.dot(W2.T) * (1 - np.power(a1, 2))
16    dW1 = np.dot(X.T, delta2)
17    db1 = np.sum(delta2, axis=0)
```

Implement

Batch gradient descent (Cont.)

Finally, here comes the function to train our Neural Network.

```
1      # Add regularization terms (b1 and b2 don't have
      regularization terms)
2      dW2 += reg_lambda * W2
3      dW1 += reg_lambda * W1
4      # Gradient descent parameter update
5      W1 += -epsilon * dW1
6      b1 += -epsilon * db1
7      W2 += -epsilon * dW2
8      b2 += -epsilon * db2
```


Implement

Batch gradient descent (Cont.)

Finally, here comes the function to train our Neural Network.

```
1  # Assign new parameters to the model
2  model = { 'W1': W1, 'b1': b1, 'W2': W2, 'b2': b2 }
3
4  # Optionally print the loss.
5  # This is expensive because it uses the whole
   dataset, so we don't want to do it too often.
6  if print_loss and i % 1000 == 0:
7      print "Loss after iteration %i: %f" %(i,
       calculate_loss(model))
8
9  return model
```

A network with a hidden layer of size 3

Let's see what happens if we train a network with a hidden layer size of 3.

```
1 # Build a model with a 3-dimensional hidden layer
2 model = build_model(3, print_loss=True)
3
4 # Plot the decision boundary
5 plot_decision_boundary(lambda x: predict(model, x))
6 plt.title("Decision Boundary for hidden layer size 3")
```

Varying the hidden layer size

Let's now get a sense of how varying the hidden layer size affects the result.

```
1 plt.figure(figsize=(16, 32))
2 hidden_layer_dimensions = [1, 2, 3, 4, 5, 20, 50]
3 for i, nn_hdim in enumerate(hidden_layer_dimensions):
4     plt.subplot(5, 2, i+1)
5     plt.title('Hidden Layer size %d' % nn_hdim)
6     model = build_model(nn_hdim)
7     plot_decision_boundary(lambda x: predict(model, x))
8 plt.show()
```